

Real-Time Heart Beat Anomaly Detection

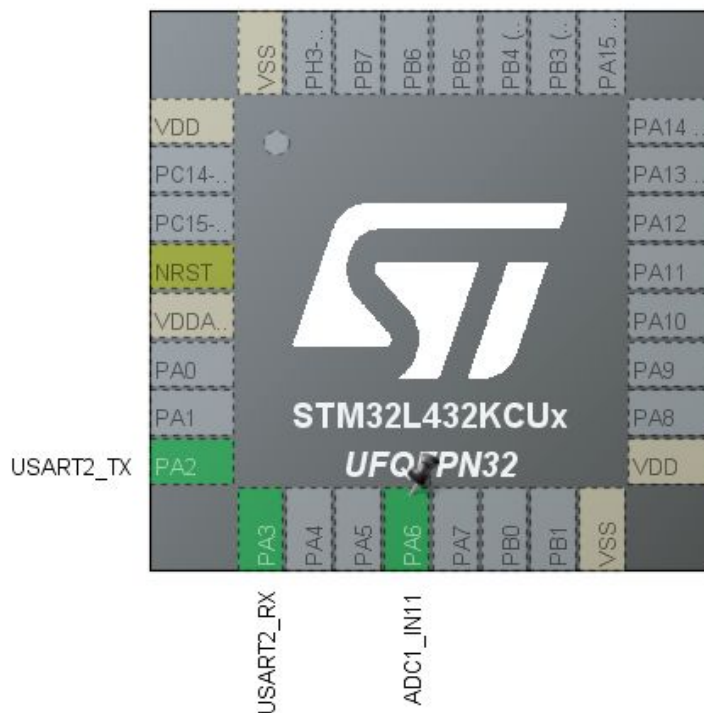
Submitted to:
Dr. Mohamed Shalan

Submitted by:
Amro Ghoneim 900150924
Ismail ElSharkawy 900151594
Zeyad Ali 900152807

Project Design and Description

The aim of our project is to detect anomalies in heartbeats constructed from an electrocardiogram (ECG) monitor, which is the recording of the electrical pulse/activity of one's heart on the MCU in real time. The recording of the heart beat will be displayed and the system will detect if there is an anomaly in the heart beat.

First the AD8232 sensor will be used to output the corresponding ADC signals. These signals will be received by the STM32L432 Nucleo microcontroller. Pin PA6 on the board was configured as shown in the following figure as an ADC input. This pin will



receive the ADC signal and the signals will then be used later for anomaly detection.

The final design of our system is as follows:

- The heart pulses will be detected using the AD8232 module.
- The readings detected from the AD8232 will be sent to the STM32 microcontroller on one of its ADC input pins.
- Needed Preprocessing tasks will be applied to the incoming readings

- The readings will then be given to the machine learning model deployed on the MCU for real time inference
- The readings/output will then be transmitted via UART and/or displayed on MCU(LEDs)

Data set

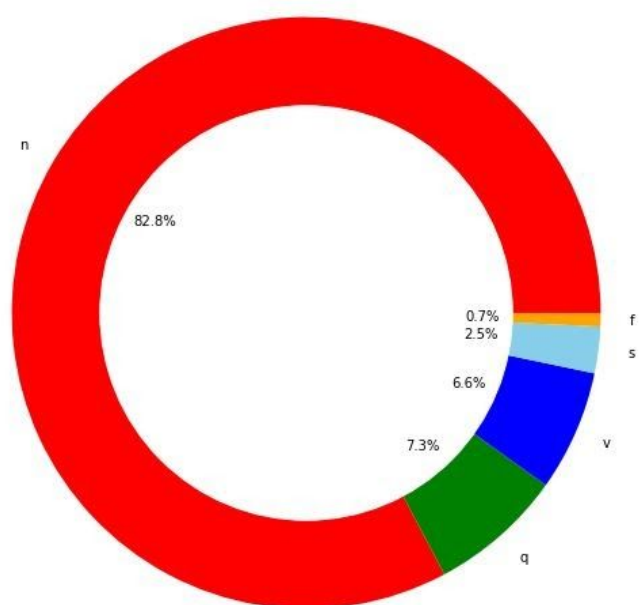
For the purposes of this project the data set used was the ECG heartbeat Categorization Data set.

It consists of two collections of heartbeat signals that are derived from two famous data sets in this domain, The MIT-BIH Arrhythmia Data set and the PTB Diagnostic Data set. For this project, we used the collection based on the MIT-BIH Data set as the number of samples are approximately 110K using 11 bit resolution and downsampled at 125 Hz. This is a good amount of samples to train a Neural Network properly. This Data set is comprised of 5 classes which are the following:

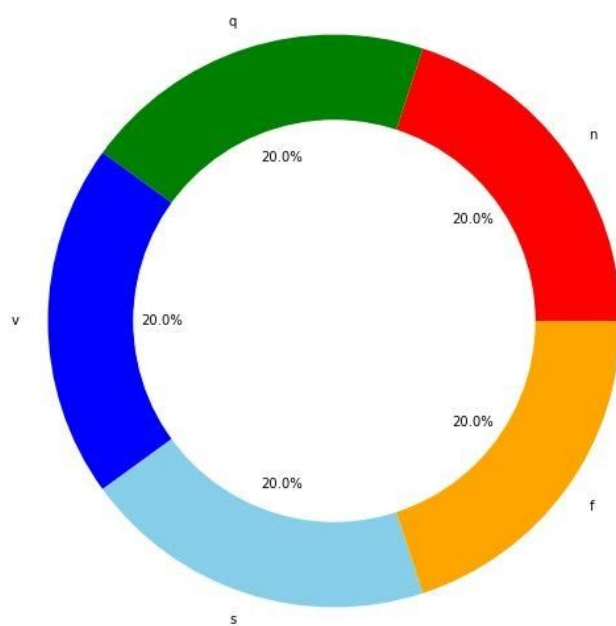
- 'N' : 0
- 'S' : 1
- 'V' : 2
- 'F' : 3
- 'Q' : 4

Where N represents the normal class and all other classes are abnormalities. The following screenshots show the class distributions in the training data before and after data augmentation:

Before data augmentation:



After Data Augmentation:



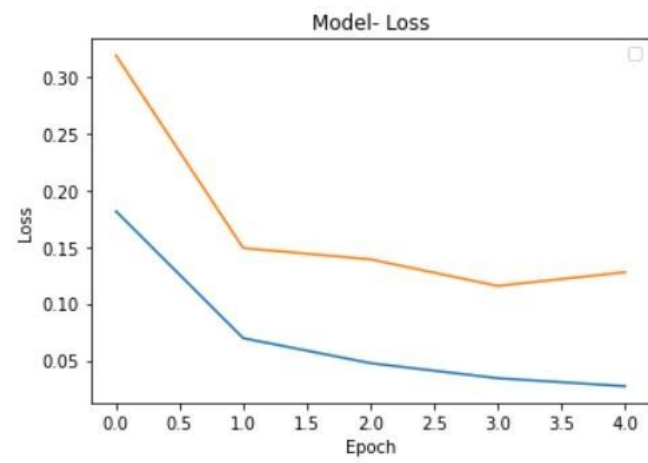
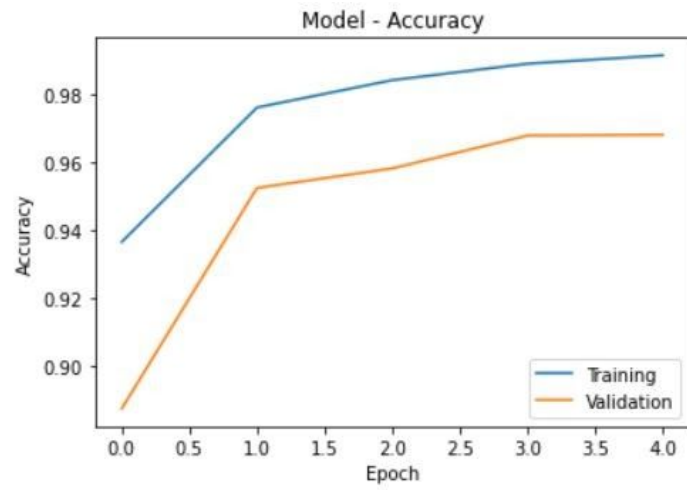
Where data augmentation restructured the data to make it more balanced, which in turn allows for better training.

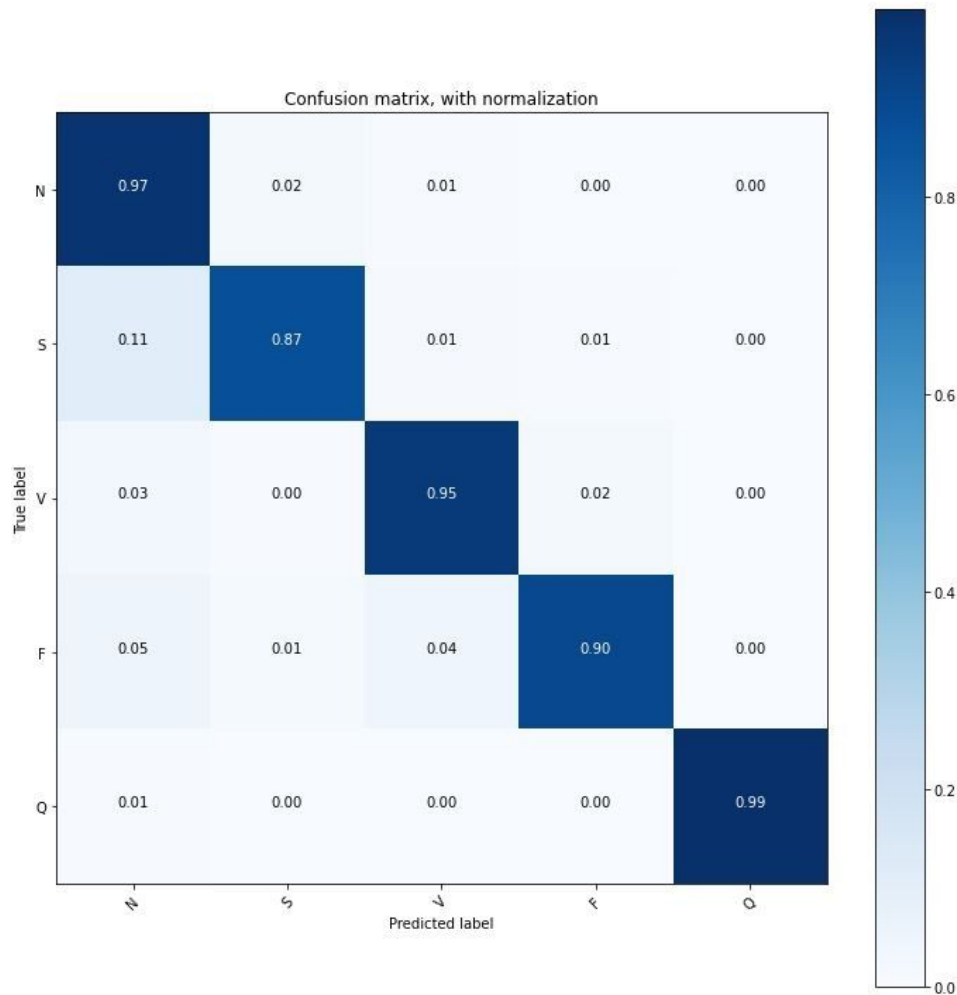
Machine Learning Model:

The initial machine learning model used was a 11 layer model (not including the output and input layers), consisting of 3 1-D CNN layers, 3 Max-Pooling layers and 3 Batch Normalization layers followed by 2 Dense layers (Keras Documentation provided in references) as seen in the screenshot below:

```
data_input=(X_train.shape[1],1)
inputs_cnn=Input(shape=(im_shape), name='data_input')
conv1_1=Convolution1D(64, (6), activation='relu', input_shape=im_shape)(inputs_cnn)
conv1_1=BatchNormalization()(conv1_1)
pool1=MaxPool1D(pool_size=(3), strides=(2), padding="same")(conv1_1)
conv2_1=Convolution1D(64, (3), activation='relu', input_shape=im_shape)(pool1)
conv2_1=BatchNormalization()(conv2_1)
pool2=MaxPool1D(pool_size=(2), strides=(2), padding="same")(conv2_1)
conv3_1=Convolution1D(64, (3), activation='relu', input_shape=im_shape)(pool2)
conv3_1=BatchNormalization()(conv3_1)
pool3=MaxPool1D(pool_size=(2), strides=(2), padding="same")(conv3_1)
flatten=Flatten()(pool3)
dense_end1 = Dense(64, activation='relu')(flatten)
dense_end2 = Dense(32, activation='relu')(dense_end1)
main_output = Dense(5, activation='softmax', name='main_output')(dense_end2)
```

The model, created using **python** and **keras** performed quite well after only 5 epochs of training as described by the screenshots below:





However, when using the model in tandem with STM32Cube.AI and the provided Nucleo MCU, it can be seen that the model, even after compression, is too large and complex to be used.

Therefore; we decided to implement a new model that is much smaller and less complex that still proved to be quite accurate.

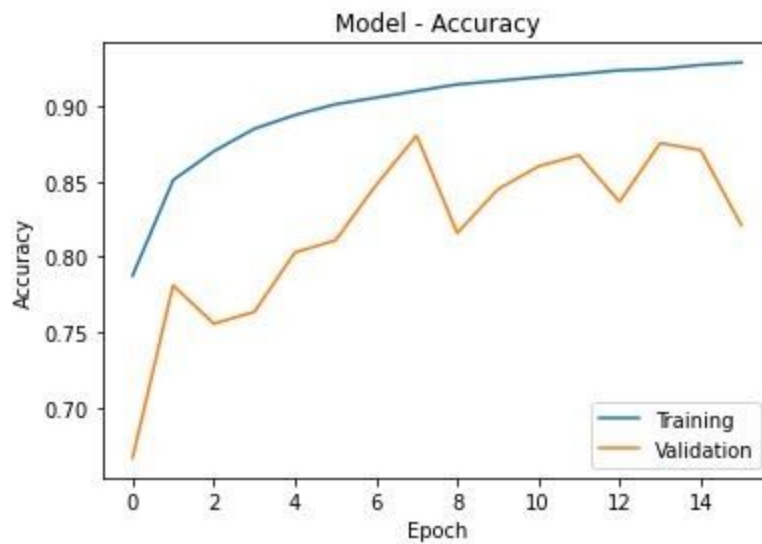
The new model architecture consists of only two dense layers (not including input and output layers). The screenshots below show the implementation of the machine learning model using **python** and **Keras** as well as the model's performance after 20 epochs of training:

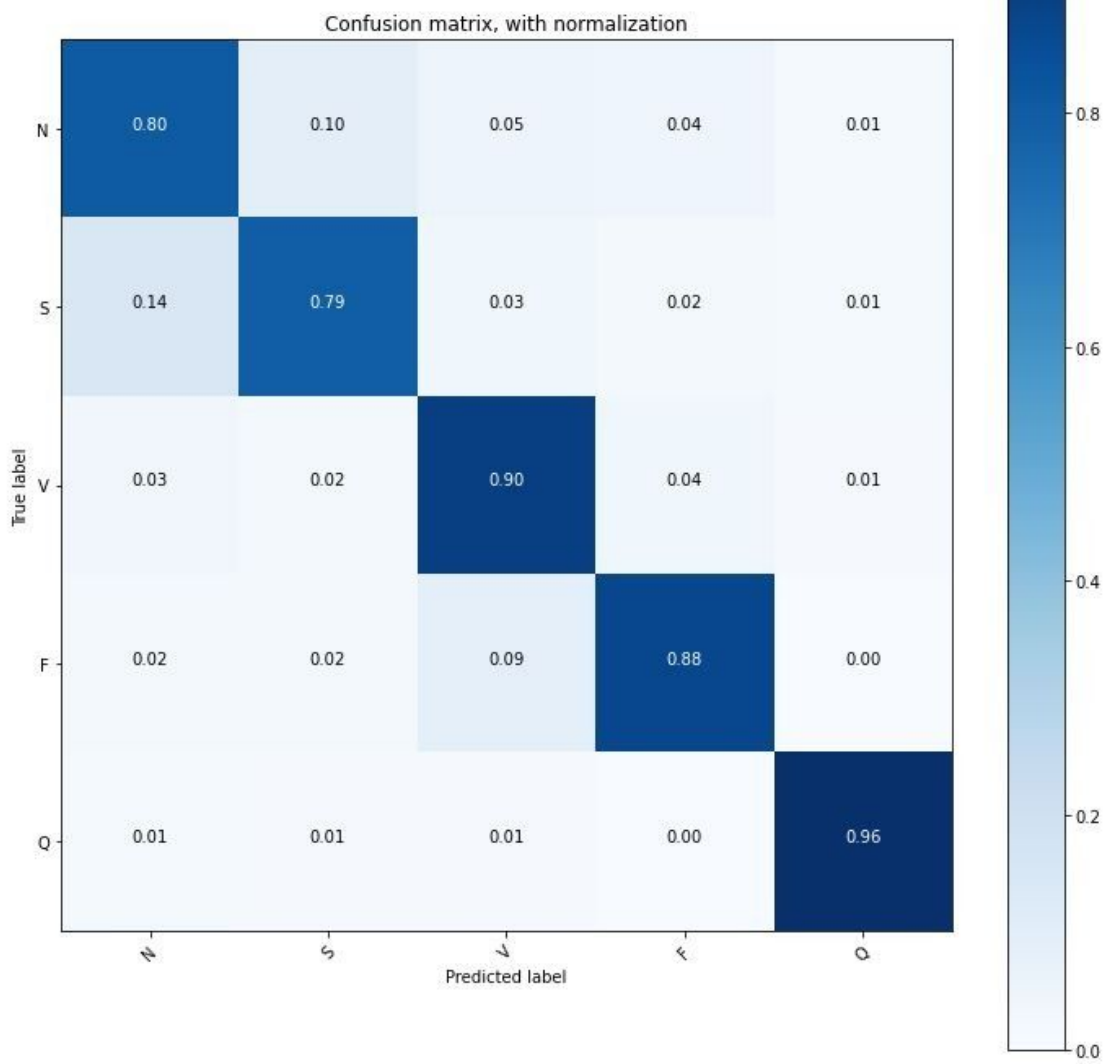
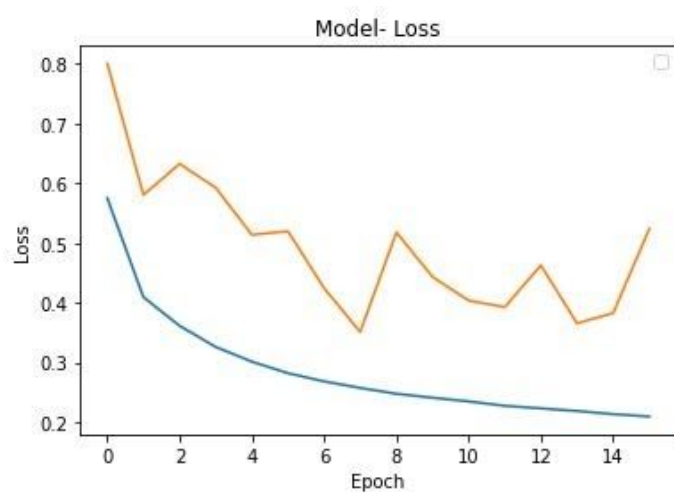
```
def network_1(X_train,y_train,X_test,y_test):

    data_input=(X_train.shape[1],1)
    inputs_dense = Input(shape = (data_input), name = 'data_input')
    dense_1 = Dense(50, activation='relu', input_shape=data_input)(inputs_dense)
    dense_2 = Dense(50, activation='relu')(dense_1)
    flat = Flatten()(dense_2)
    dense_out = Dense(5, activation='softmax', name='main_output')(flat)

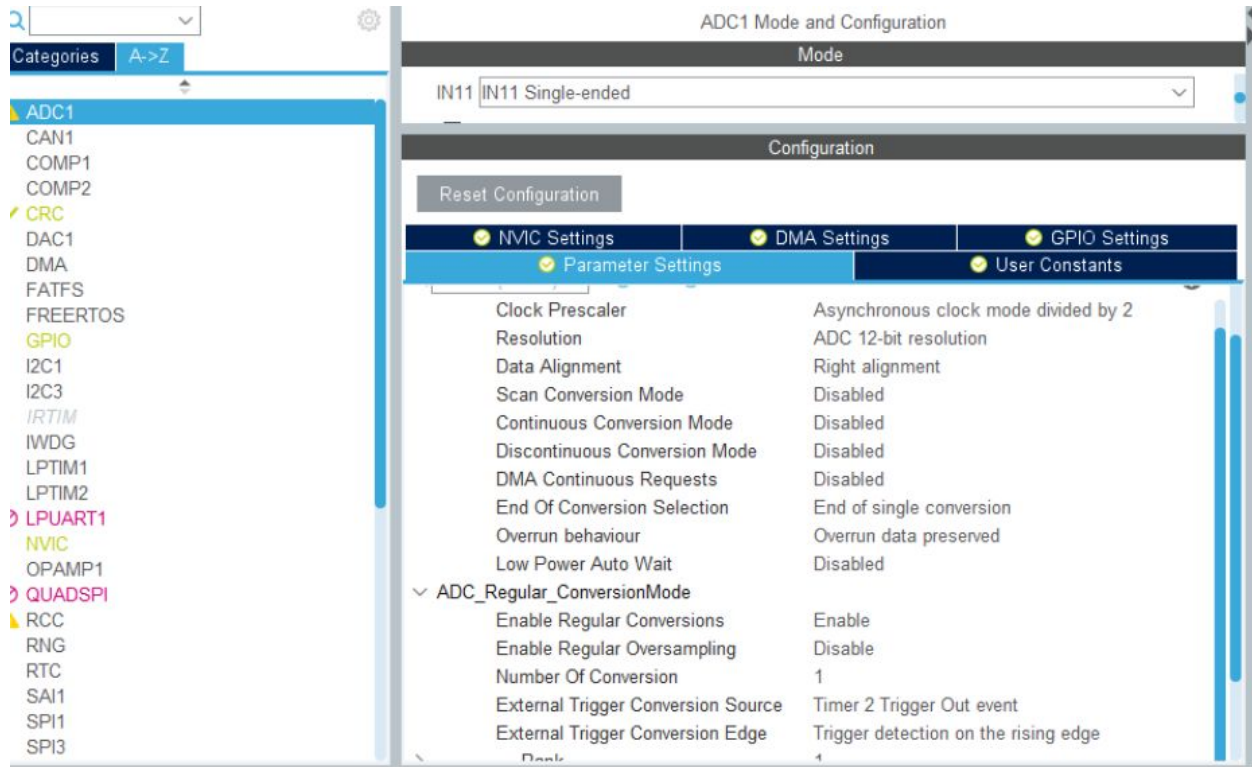
    model = Model(inputs= inputs_dense, outputs=dense_out)

    model.compile(optimizer='adam', loss='categorical_crossentropy',metrics = ['accuracy'])
```





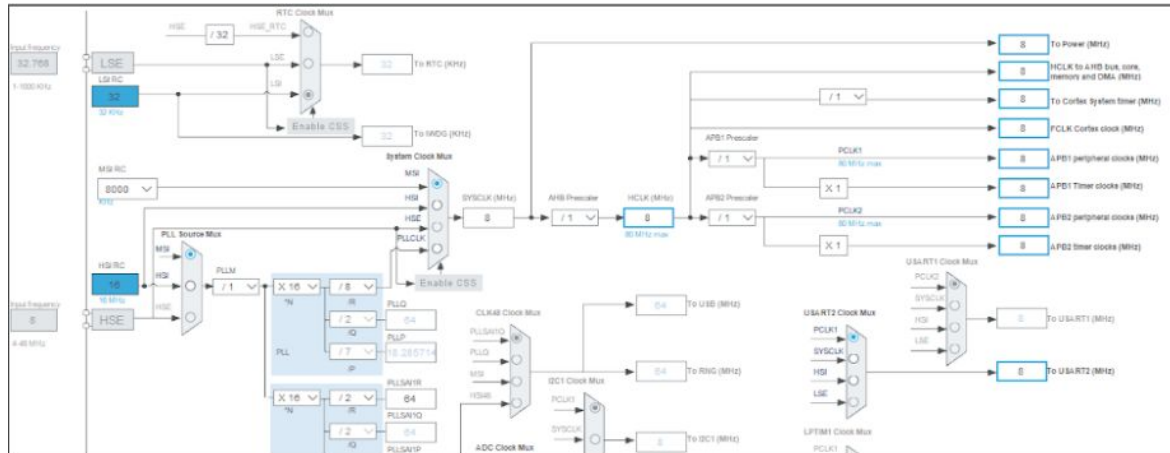
CubeMX settings for ADC and Timer:



Here, we see the specifications of our used ADC. As we can see, we use a 12 bit resolution with enabled regular conversions alongside a Timer 2 trigger out event on the rising edge of our clock. (1 conversion per trigger)

✔ NVIC Settings		✔ DMA Settings		✔ GPIO Settings	
✔ Parameter Settings			✔ User Constants		
NVIC Interrupt Table		Enabled	Preemption Priority		Sub Priority
ADC1 global interrupt		<input checked="" type="checkbox"/>	0		0

We ensure that the global interrupt is enabled so we can follow desired logic each time a conversion is made.



Here we see our clock tree, the one relevant to us is the APB1 Timer clock which ends up initially being 8MHz.

Reset Configuration

Parameter Settings

User Constants

NVIC Settings

DMA Settings

Configure the below parameters :

Counter Settings

Prescaler (PSC - 16 bits value)

7

Counter Mode

Up

Counter Period (AutoReload Register...)

7999

Internal Clock Division (CKD)

No Division

auto-reload preload

Disable

Trigger Output (TRGO) Parameters

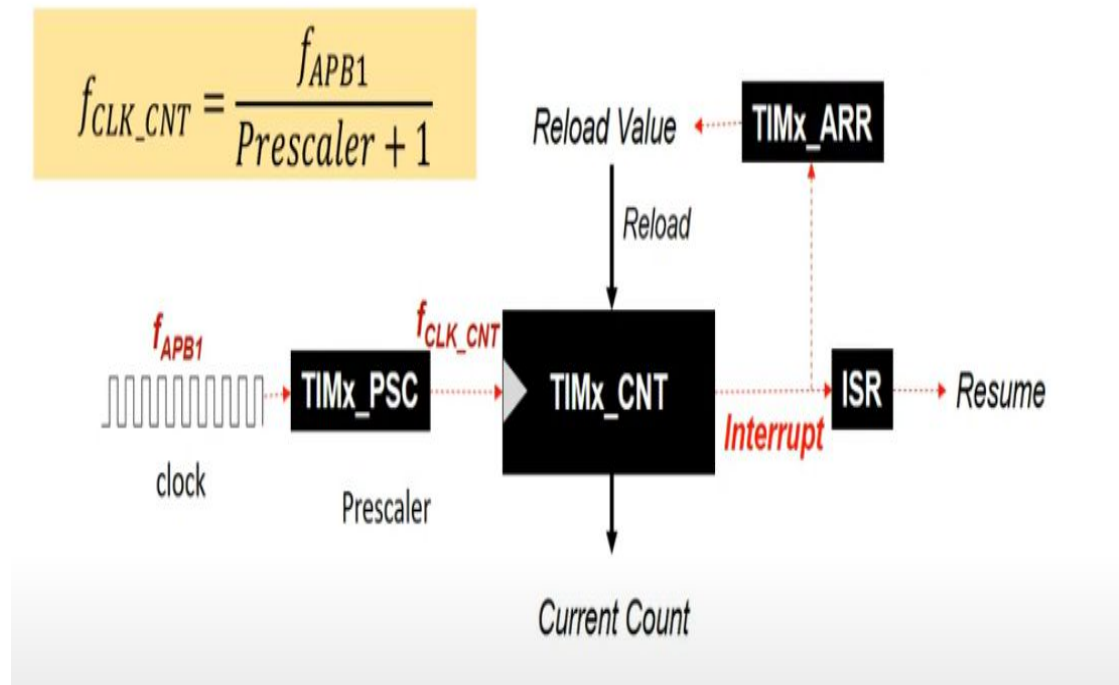
Master/Slave Mode (MSM bit)

Disable (Trigger input effect not delayed)

Trigger Event Selection TRGO

Update Event

As seen in the figure above, our prescaler is set to 7, and out counter period goes up to 7999.



From this figure, we see how the sampling process took place according to previously mentioned assigned values. $f_{APB1} = 8\text{MHz}$; Since the prescaler is set to 7, the f_{CLK_CNT} is 1MHz. That 1MHz used to calculate the interrupt frequency by being divided by the counter period value + 1 which results in a frequency of 125 Samples/Sec or 1 sample every 8ms.

CubeMX settings FOR AI:

The screenshot below shows a summary of the model's size on chip as well as the model used with the compression rate specified. This initial step acts as a filter for MCUs where any MCU that cannot include AI or this specific model size will not be shown

New Project

MCU/MPU Selector | Board Selector | Cross Selector

23 168

Eprom = 0 (Bytes)

Flash From 16 to 2048 (kBytes)

16 2048

Ram From 16 to 1184 (kBytes)

16 1184

Freq. From 48 to 480 (MHz)

48 480

Artificial Intelligence

☒ Enable

Model Keras

Type Saved model

Model my_model.h5

Browse

Compression 8

Analyze

STM32Cube embedded software now available on GitHub

AI Summary

Keras Minimum Flash: 6.69 KBytes Minimum Ram: 400.00 KBytes

C:\Users\world\Downloads\my_model.h5

MCUs/MPUs List: 908 items

Part No.	Reference	Marketing Status	Unit Price for 1000 (USD)	Board	Package	Flash	RAM	ID	Freq
STM32F301C6	STM32F301C6Tx	Active	1.596		LQFP48	32 kBytes	16 kBytes	37	72 MHz
STM32F301C8	STM32F301C8Tx	Active	1.666		LQFP48	64 kBytes	16 kBytes	37	72 MHz
STM32F301C8Yx	STM32F301C8Yx	Active	1.666		WLCSF49	64 kBytes	16 kBytes	37	72 MHz
STM32F301K6	STM32F301K6Tx	Active	1.272		LQFP32	32 kBytes	16 kBytes	25	72 MHz
STM32F301K6Jx	STM32F301K6Jx	Active	1.272		UFQFPN32	32 kBytes	16 kBytes	24	72 MHz
STM32F301K8	STM32F301K8Tx	Active	1.342		LQFP32	64 kBytes	16 kBytes	25	72 MHz
STM32F301K8Jx	STM32F301K8Jx	Active	1.342		UFQFPN32	64 kBytes	16 kBytes	24	72 MHz
STM32F301R6	STM32F301R6Tx	Active	1.758		LQFP64	32 kBytes	16 kBytes	51	72 MHz
STM32F301R8	STM32F301R8Tx	Active	1.828		LQFP64	64 kBytes	16 kBytes	51	72 MHz
STM32F302C6	STM32F302C6Tx	Active	1.712		LQFP48	32 kBytes	16 kBytes	37	72 MHz
STM32F302C8	STM32F302C8Tx	Active	1.782		LQFP48	64 kBytes	16 kBytes	37	72 MHz
STM32F302C8Yx	STM32F302C8Yx	Active	1.782		WLCSF49	64 kBytes	16 kBytes	37	72 MHz
STM32F302CB	STM32F302CBTx	Active	1.99		LQFP48	128 kBytes	32 kBytes	37	72 MHz
STM32F302CC	STM32F302CCTx	Active	2.288		LQFP48	256 kBytes	40 kBytes	37	72 MHz
STM32F302K6	STM32F302K6Jx	Active	1.596		UFQFPN32	32 kBytes	16 kBytes	24	72 MHz
STM32F302K8	STM32F302K8Jx	Active	1.666		UFQFPN32	64 kBytes	16 kBytes	24	72 MHz
STM32F302R6	STM32F302R6Tx	Active	1.974		LQFP64	32 kBytes	16 kBytes	51	72 MHz

After choosing the nucleo board, additional settings must be set as shown below:

Additional Software Components selection

Filters

Search

Pack Vendor

Software Component Class

Packs

Pack / Bundle / Component	Version	Selection
ARM.CMSIS	5.7.0	
STMicroelectronics X-CUBE-AI	5.0.0	
Artificial Intelligence X-CUBE-AI	5.0.0	
Core		<input checked="" type="checkbox"/>
Artificial Intelligence Application	5.0.0	
Application		ApplicationTemplate
STMicroelectronics X-CUBE-ALGOBUILD	1.0.0	
STMicroelectronics X-CUBE-BLE1	5.0.0	
STMicroelectronics X-CUBE-BLE2	1.0.0	
STMicroelectronics X-CUBE-GNSS1	4.1.0	
STMicroelectronics X-CUBE-MEMS1	7.2.0	
STMicroelectronics X-CUBE-NFC4	1.5.2	
STMicroelectronics X-CUBE-SUBG2	1.1.0	
STMicroelectronics X-CUBE-TOUCHGFX	4.13.0	

Component dependencies

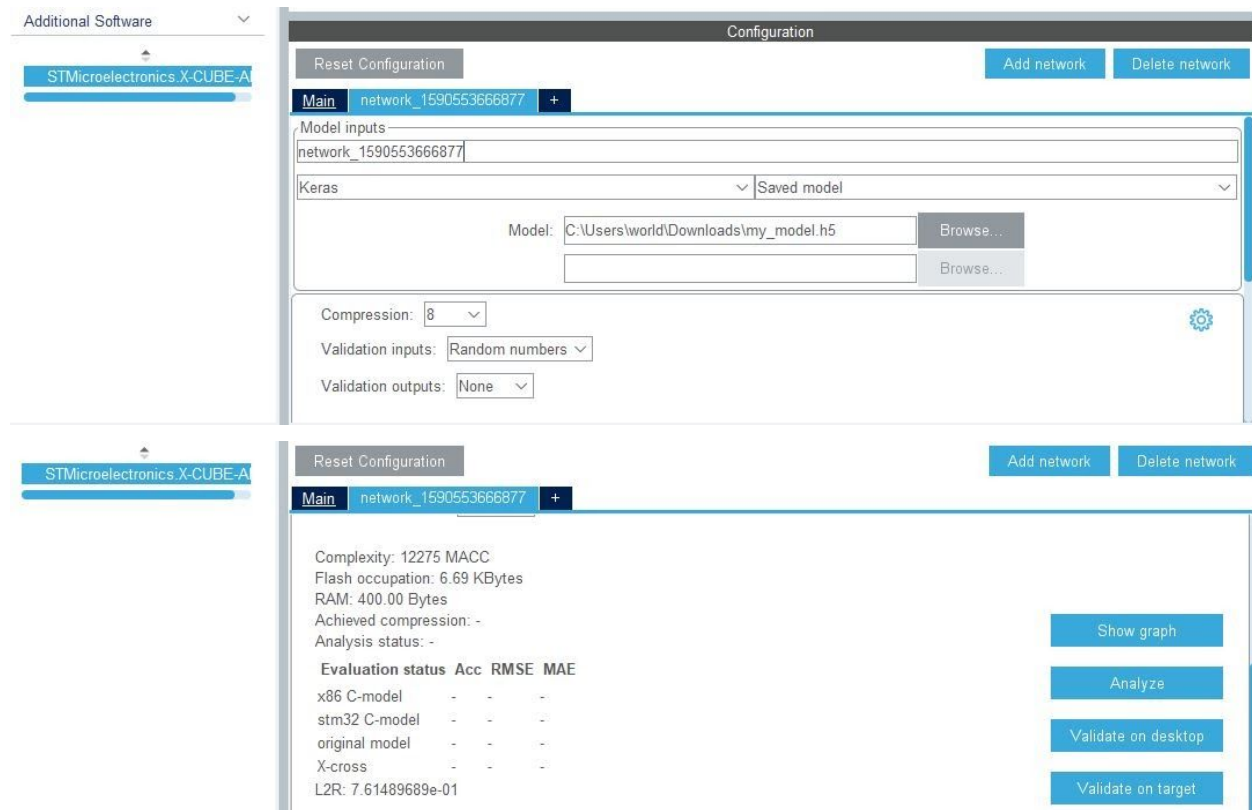
Details and warnings

Click on a pack / bundle / component to see the details

Show/hide filters Show/hide details Show/hide dependencies

Ok Cancel

The screenshots below show the network as included before with a random name. This name should be changed as it will reflect in the generated code. They also show some statistics and methods that can be used to see how the model fares after compression.



Implementation(Keil/main):

```
void HAL_ADC_ConvCpltCallback(ADC_HandleTypeDef *hadc)
{
    /* Prevent unused argument(s) compilation warning */
    UNUSED(hadc);
    /* NOTE : This function should not be modified. When the callback is needed,
       function HAL_ADC_ConvCpltCallback must be implemented in the user file.
    */
    HAL_ADC_PollForConversion(&hadc1, HAL_MAX_DELAY);
    adcval = HAL_ADC_GetValue(&hadc1); //getting the ADC value
    //adcval= ((float)adcval/4096.0) * 2048.0;
    if(count<187){ //only add the first 187 samples to the buffer in_data
        in_data[count] = adcval;
        if(in_data[count]<min) //get the minimum value
            min=in_data[count];
        if(in_data[count]>max) //get the maximum value
            max=in_data[count];
        count++;
    }
    sprintf(reading, "%hu", adcval); // copy the adc value to a reading buffer to display the readings if needed.
    HAL_UART_Transmit(&huart2, (uint8_t*)reading, sizeof(reading), HAL_MAX_DELAY); //transmit reading
    HAL_UART_Transmit(&huart2, (uint8_t*)newline, sizeof(newline), HAL_MAX_DELAY); //newline
}
```


In the beginning, we receive conversions from the ADC and store them in *adcval* once every 8ms so that we can reach a total of 125 samples per second (like reference paper). Our received converted values are then stored in the *in_data buffer* before we proceed to find maximum and minimum of each 187 values interval and incrementing the counter.

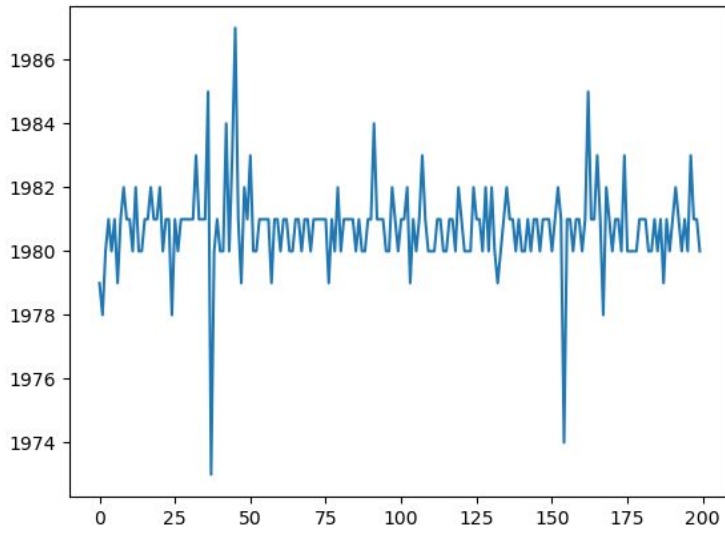
```
while (1)
{
    if(count>=187){ //if Counter exceeded 187 samples
        HAL_ADC_Stop_IT(&hadcl); //stop ADC interrupts so no more values are read
        for (int j=0; j<187;j++){
            norm_in_data[j]= (in_data[j]-min)/(max-min); //normalize the signals in in_data buffer
        }
        test = MX_X_CUBE_AI_Process(norm_in_data); //use the AI_Process to get the prediction
        if(!isnan(test))
            {sprintf(inference, "%f\r\n", test); //copy the probability to inference array
            //HAL_UART_Transmit(&huart2, (uint8_t*)inference, sizeof(inference), HAL_MAX_DELAY);
            if((inference[2]=='9') || (inference[0]=='1')) //check if the probability is >=0.9
                HAL_UART_Transmit(&huart2, (uint8_t*)normal, sizeof(normal), HAL_MAX_DELAY); //transmit normal
            else
                HAL_UART_Transmit(&huart2, (uint8_t*)abnormal, sizeof(abnormal), HAL_MAX_DELAY); //transmit abnormal
            }
        count =0; //reset counter
        min=999999; //reset minimum
        max=-1; //reset maximum
        HAL_ADC_Start_IT(&hadcl); //start ADC interrupt again.
    }
}
```

After the collection of 187 samples, interrupts are stopped and we proceed to normalize all collected samples before they are sent to the *MX_X_CUBE_AI_Process* function which returns the prediction. We make sure the returned value isn't null and transmit prediction as either *Normal* or *Abnormal* depending on a threshold of 0.9 plus. At the end of any 187 value interval, values get reset and the interrupts are started again.

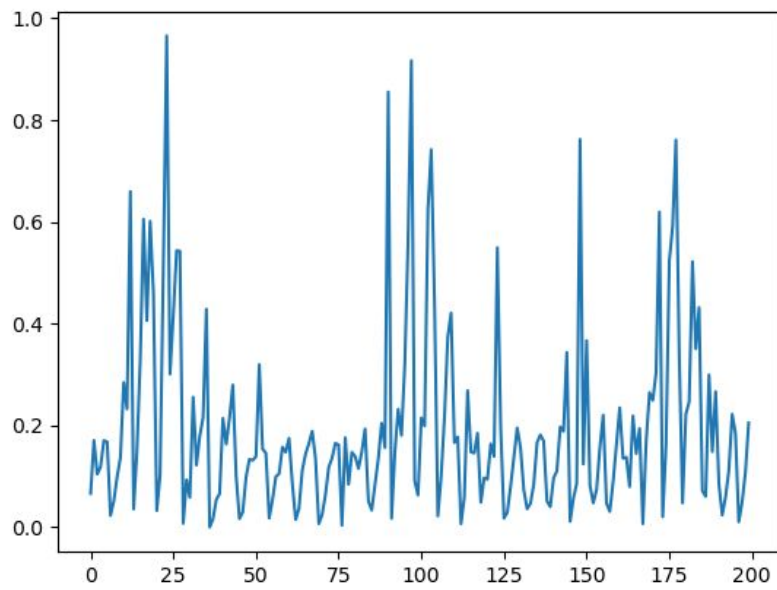
Results:

The screenshots below show the output signal as well as the normalized output signal that is fed into the machine learning model for around 200 samples.

ADC Signal:

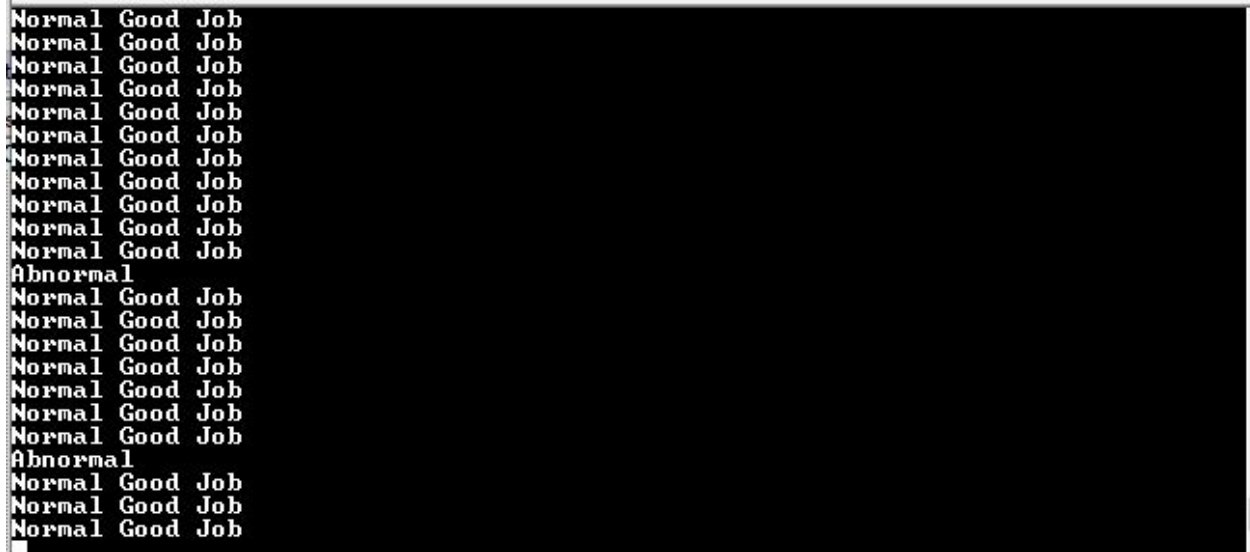


Normalized ADC Signal:



The screenshot below shows a sample of the output predictions in real time after feeding the normalized data into the model.

Predictions:



```
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Abnormal
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Normal Good Job
Abnormal
Normal Good Job
Normal Good Job
Normal Good Job
```

Comments on the results:

We were able to gather the data from the ECG sensor in accordance with the required sampling rate (125 Hz) and the sampling resolution (11 bits). We were able to do some preprocessing tasks such as normalizing the data before feeding it into the machine learning model. The model performs the inference and returns the results to be displayed. The results achieved are not quite accurate, however, due to a number of limitations. The first one is the data coming from the sensor is quite noisy as can be seen in the screenshots provided earlier. In addition, memory limitations regarding the nucleo MCU and the Keil IDE led to the usage of a highly compressed model with somewhat degraded results as well as not being able to complete the required

preprocessing tasks that are described by the Data set's paper (provided in the references). Therefore, further work regarding the preprocessing of the data is required to allow for better results.

References:

1. <https://www.physionet.org/physiobank/database/mitdb/>
2. <https://keras.io/api/>
3. <https://www.kaggle.com/shayanfazeli/heartbeat>
4. https://www.st.com/resource/en/user_manual/dm00570145-getting-started-with-x-cubeai-expansion-package-for-artificial-intelligence-ai-stmicroelectronics.pdf
5. https://www.st.com/content/st_com/en/stm32-ann.html